

# Decision Trees for Classification

Zhu Pengxiang\*, Zhou Shengyang<sup>†</sup>

2023 Spring

## Abstract

This paper aims to discuss one of the important applications of information theory in machine learning, i.e. decision trees. We first study the formation of the original ID3 algorithm and the subsequent C4.5 and CART algorithm. Through the process we have derived an intuitive understanding of how entropy can be used for filtering out attributes for better classification. Then we conduct some empirical experiments on a dataset to compare the performance of the three algorithms we have introduced. We have found that the three algorithms have *similar performances* on the dataset. For the final part, we utilize the knowledge obtained in lectures to propose a modified version of decision tree algorithm using a novel pruning technology, inspired by the *dropout* method in neural network. We have found that under certain hyper-parameter settings, our modified algorithm can *achieve better performance* than the three classic algorithms when the test set contains new combinations of data inputs and high probability of overfitting, but with less stability and inferior performance when the training set is well-defined and comprehensive. The code of the modified algorithm that we have proposed can be found in <https://github.com/JubSteven/ICE2601-Group-Project>.

---

\*Responsible for writing fundamental knowledge, ID3 algorithm and empirical studies. Also responsible for all the code required for the paper. The proposed modification and innovation is formulated under mutual discussion

<sup>†</sup>Responsible for writing the analysis of classic models and pruning technology. Also does the slides. The proposed modification and innovation is formulated under mutual discussion.

# Contents

<b>1</b>	<b>Basic Concepts</b>	<b>1</b>
1.1	Fundamentals of Information Theory . . . . .	1
1.2	General Formation of Classification Problems . . . . .	1
<b>2</b>	<b>ID3 Tree Induction Algorithm</b>	<b>2</b>
2.1	General Setting of a Decision Tree . . . . .	2
2.2	Information Gain as Selection Measure . . . . .	3
<b>3</b>	<b>Variations of Decision Tree Algorithm</b>	<b>4</b>
3.1	C4.5 Algorithm . . . . .	4
3.2	CART Algorithm . . . . .	5
3.3	Tree Pruning . . . . .	6
<b>4</b>	<b>A Modified Decision Tree Algorithm</b>	<b>7</b>
4.1	Results of Classic Algorithms . . . . .	7
4.2	Modified Pruning based on Dropout . . . . .	9
4.3	Strengths and Limitations . . . . .	11

# 1 Basic Concepts

## 1.1 Fundamentals of Information Theory

Entropy, proposed by Shannon in [6], is a measure of the uncertainty of a random variable. Consider a discrete random variable  $X \sim p(x)$  with alphabet  $\mathcal{X}$ , the entropy  $H(X)$  is defined by

$$H(X) = - \sum_{p(x)} p(x) \log p(x) \quad (1)$$

Suppose two random variables  $X, Y \sim p(x, y)$  with alphabets  $\mathcal{X}, \mathcal{Y}$ , the conditional entropy  $H(Y|X)$  is defined by

$$H(Y|X) = \sum_{x \in \mathcal{X}} \sum_{y \in \mathcal{Y}} p(x, y) \log p(y|x) \quad (2)$$

As stated in [2], the mutual information  $I(X; Y)$  is a measure of the reduction in the uncertainty of one random variable due to the knowledge of the other. It can be regarded as a measure of information gain. It satisfies

$$I(X; Y) = H(X) - H(X|Y) \quad (3)$$

## 1.2 General Formation of Classification Problems

Data classification problems consist of two parts, the learning step and classification step [3]. During the learning step, a training data set  $D$  of length  $N$  needs to be prepared. Here we suppose that there are  $n$  attributes for each sample  $\mathbf{X}_i \in D$ ,  $i = 1, 2, \dots, N$ . For each potential  $\mathbf{X}$  in alphabet  $\mathcal{X}$ , we assume that there are  $k$  possible labels, denoted by  $C = \{c_1, \dots, c_k\}$ . A predefined label  $c(\mathbf{x}) \in C$  is associated with each sample  $\mathbf{X}_i$ . The training process is intended to formulate a function  $f$  such that

$$f := \mathcal{X} \mapsto C \quad (4)$$

Nowadays, numerous methods are proposed to train the mapping from a given training dataset. Decision trees along with its numerous variants, is a typical approach to this problem. In this paper, we will introduce the basic algorithm and certain optimizations. We will then focus on its performance and explain from the perspective of information theory the limitations of the classic ID3 algorithm.

## 2 ID3 Tree Induction Algorithm

### 2.1 General Setting of a Decision Tree

Decision tree induction is the learning of decision trees of class-labeled training tuples. A typical decision tree is shown in Figure 1 (from our work in CUMCM 2022, used CART algorithm). It consists of a tree structure where each branch represents an outcome of the decision process and each node has a label that denotes the splitting criterion. Given a test data  $\mathbf{X}_t$  where the associated label is unknown, then we can construct a path from the root node to one of the leaf node, which corresponds to one of the labels. It is then returned as the final result.

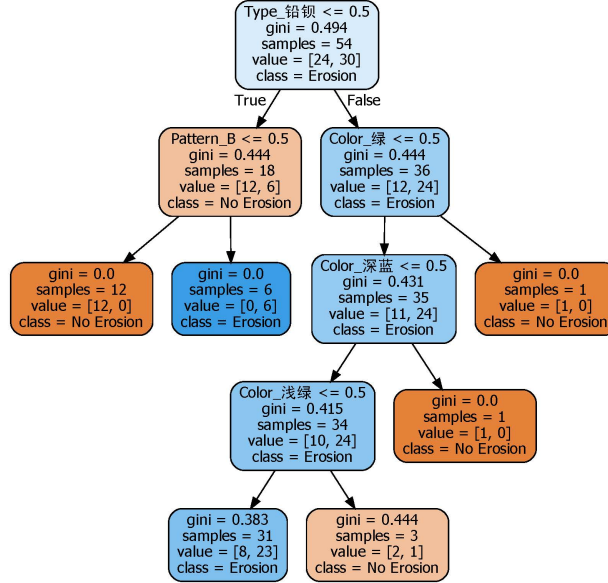


Figure 1: A typical example of decision tree

Now we introduce the ID3 algorithm, first proposed by [4], to construct such a tree from training samples. The algorithm is illustrated as follows[3]. The inputs are a data partition  $D$  consisting of training tuples and class labels. Since this is a recursive algorithm in a top-down approach,  $D$  is the complete training set. `attribute_list` refers to the list of attributes ( $n$  of them in total) used for classification, and `attribute_selection_method` is a method to determine the splitting criterion. Here we consider the case of ID3 algorithm, and it will alter in the subsequent optimizations. For simplicity, we only consider the cases where all attributes are discrete.

---

**Algorithm 1:** Generate\_Decision\_Tree

---

**Input:**  $D$ , attribute\_list, attribute\_selection\_method

**Output:** A decision tree  $T$

create a node  $N$ ;

**if** for all  $\mathbf{x} \in D$ ,  $c(\mathbf{x})$  are the same **then**

└ return  $N$  as a leaf node labeled with  $c(\mathbf{x})$

**if** attribute\_list is empty **then**

└ return  $N$  as a leaf node labeled with the majority class in  $D$ ;

apply attribute\_selection\_method( $D$ , attribute\_list) to find the best splitting criterion;

label node  $N$  with splitting criterion;

attribute\_list = attribute\_list - splitting\_attribute;

**for** each outcome  $j$  of splitting\_criterion **do**

└ Denote  $D_j \subseteq D$  such that  $D_j$  satisfies outcome  $j$ ;

└ **if**  $D_j = \emptyset$  **then**

└└ attach a leaf labeled with the majority  $D$  to node  $N$ ;

└ **else**

└└ attach the node returned by Generate\_Decision\_Tree( $D_j$ , attribute\_list) to node  $N$ ;

return  $N$ ;

---

Given a training set  $D$ , the time complexity required for the algorithm is  $O(n|D| \log |D|)$ .

From the algorithm, we can see that finding the right splitting criterion is crucial for the structure of the tree, which further influences the test accuracy as well as robustness of the classifier. We will further our discussion in the following part.

## 2.2 Information Gain as Selection Measure

The ID3 algorithm utilizes information gain, which is essentially mutual information, as attribute selection measure. Suppose node  $N$  consists of a data partition  $D$ , then the attribute with highest gain is selected. By this approach, the information needed to classify new resulting partitions  $D_j$  is minimized. We will formalize the problem as follows.

Suppose we want to partition tuples in  $D$  on some attribute  $A$  having  $v$  distinct values  $\{a_1, \dots, a_v\}$ . Since we are considering the discrete case, we can split  $D$  into  $v$  partitions such that  $D_j \subseteq D$  satisfies  $\mathbf{X} \in D_j$  and  $\mathbf{X}_A = a_j$ . Ideally, we intend to make each partition pure, i.e all elements in  $D_j$  belonging to the same class. But in many cases it is impossible to achieve the goal, so we introduce entropy to measure the amount of information we still need for exact classification.

From information theory we know that the expected information needed to classify a tuple in  $D$  is simply its entropy

$$H(D) = - \sum_{i=1}^m p_i \log_2(p_i) \quad (5)$$

where  $p_i$  is the probability that  $\mathbf{X} \in D$  belongs to class  $c_i$ , so it can be estimated by  $\frac{|c_i|}{|D|}$ . The conditional entropy  $H(D|A)$  can also be derived using the following equation.

$$H(D|A) = \sum_{j=1}^v p(A = v_j) H(D|A = v_j) = \sum_{j=1}^v \frac{|D_j|}{|D|} H(D_j) \quad (6)$$

Thus, we have derived the formula for information gain, which is an effective measure on how much would be gained by branching on  $A$ .

$$I(D; A) = H(D) - H(D|A) \quad (7)$$

In ID3 algorithm, we chose the `splitting_attribute` as  $\arg \max_A I(D; A)$ . Since the number of branches is not limited, we can simply construct  $v$  branches from node  $N$ , with each corresponding to a partition  $D_j$ . For continuous cases, we typically sort the values of an attribute  $A$ , supposing that there are  $v$  of them. Then we can construct  $v - 1$  split points, with each split point being the mid-point between each adjacent pair.

## 3 Variations of Decision Tree Algorithm

### 3.1 C4.5 Algorithm

The information gain measure is biased towards tests with many outcomes. That is, it prefers to select attributes having a large number of values. For example, consider an attribute that acts as a unique identifier such as product ID. A split on product ID would result in a large number of partitions, in worst cases with each one containing just one tuple. Because all the partitions are pure, the information required to classify data set  $D$  based on this partitioning would be  $\text{Info}_{\text{product.ID}}(D) = 0$ . Therefore, the information gained by partitioning on this attribute is maximal. However, such a partitioning is useless for classification.

In order to solve this problem, C4.5 algorithm [5] proposes a new partition standard called Gain Ratio. In addition to using classical information entropy, the algorithm defines an interesting "split information" to describe the attribute.

Specifically, we first still calculate  $I(A; D)$ . Instead of using  $I(A; D)$  as the only indicator of reference, we introduce  $\text{SplitInfo}_A(D)$  to describe the instability of the attribute.  $\text{SplitInfo}_A(D)$  is defined as:

$$\text{SplitInfo}_A(D) = - \sum_{i=1}^v \frac{|D_j|}{D} \times \log_2 \left( \frac{|D_j|}{D} \right) \quad (8)$$

It is not difficult to see the similarity between split information and entropy: split information just replace  $p_i$  in entropy with  $\frac{|D_j|}{D}$ . In C4.5 algorithm, split information is used to measure the uncertainty of  $A$ . This value represents the potential information generated by splitting the training data set,  $D$ , into  $v$  partitions, corresponding to the  $v$  outcomes of a test on attribute  $A$ . It considers the number of tuples having that outcome with respect to the total number of tuples in  $D$ . Finally, the Gain Ratio is defined as:

$$\text{GainRatio}(A) = \frac{I(A; D)}{\text{SplitInfo}_A(D)} \quad (9)$$

The attribute with the maximum gain ratio is selected as the splitting attribute. In other words, we'll tend to choose the attribute with the least uncertainty. However, as the split information approaches 0, the ratio becomes unstable. A constraint is added to avoid this, whereby the information gain of the test selected must be large enough such that the average gain over all tests can be examined.

## 3.2 CART Algorithm

ID3 and CART [1] were invented independently of one another at around the same time, yet follow a similar approach for learning decision trees from training tuples.

However, unlike the two algorithms mentioned above, this algorithm does not use entropy as a measure, but instead uses the **Gini index** as a substitute. Using the notation previously described, the Gini index is defined as

$$\text{Gini}_A(D) = 1 - \sum_{i=1}^m p_i^2 \quad (10)$$

where  $p_i$  is the probability that a tuple in  $D$  belongs to class  $C_i$  and is estimated by  $|C_{i,D}|/|D|$ . The sum is computed over  $m$  classes. In fact, CART uses a first-order approximation to simplify time-consuming logarithmic calculations:  $\log x \sim 1 - x$ . Due to the existence of approximation, we can consider the Gini index as a special kind of entropy. When considering a binary split, we

compute a weighted sum of the impurity of each resulting partition. For example, if a binary split on  $A$  partitions  $D$  into  $D_1$  and  $D_2$ , the Gini index of  $D$  given that partitioning is

$$\text{Gini}(D) = \frac{|D_1|}{|D|} \text{Gini}(D_1) + \frac{|D_2|}{|D|} \text{Gini}(D_2) \quad (11)$$

For each attribute, each of the possible binary splits is considered. For a discrete-valued attribute, the subset that gives the minimum Gini index for that attribute is selected as its splitting subset.

Like ID3, when we need to process the attribute with continuous-valued, we'll determine the best **split-point** for it. For continuous-valued attributes, each possible split-point must be considered. The strategy is similar to that described earlier for information gain, where the midpoint between each pair of (sorted) adjacent values is taken as a possible split-point. The point giving the minimum Gini index for a given (continuous-valued) attribute is taken as the split-point of that attribute.

The reduction in impurity that would be incurred by a binary split on a discrete- or continuous-valued attribute  $A$  is

$$\Delta \text{Gini}(D) = \text{Gini}(D) - \text{Gini}_A(D) \quad (12)$$

The attribute that maximizes the reduction in impurity (or, equivalently, has the minimum Gini index) will be selected as the splitting attribute.

### 3.3 Tree Pruning

When a decision tree is built, many of the branches will reflect anomalies in the training data due to noise or outliers. Tree pruning methods address this problem of overfitting the data. There are two main approaches of tree pruning : pre-pruning approach and post-pruning approach.

In the prepruning approach, a tree is “pruned” by halting its construction early (e.g., by deciding not to further split or partition the subset of training tuples at a given node). Upon halting, the node becomes a leaf. The leaf may hold the most frequent class among the subset tuples or the probability distribution of those tuples.

When constructing a tree with the approach of pre-pruning, measures such as statistical significance, information gain, Gini index, and so on, can be used to assess the goodness of a



split. If partitioning the tuples at a node would result in a split that falls below a pre-specified threshold, then further partitioning of the given subset is halted. There are difficulties, however, in choosing an appropriate threshold. High thresholds could result in the failure of split leaves, which oversimplified the decision tree. Low thresholds could result in a few simplification.

The second and more common approach is postpruning, which removes subtrees from a “fully grown” tree. A subtree at a given node is pruned by removing its branches and replacing it with a leaf. In fact, in C4.5 and CART algorithm, they both use postpruning approach.

The cost complexity pruning algorithm used in CART is an example of the postpruning approach. This approach considers the cost complexity of a tree to be a function of the number of leaves in the tree and the error rate of the tree (where the error rate is the percentage of tuples misclassified by the tree). It starts from the bottom of the tree. For each internal node,  $N$ , it computes the cost complexity of the subtree at  $N$ , and the cost complexity of the subtree at  $N$  if it were to be pruned (i.e., replaced by a leaf node). The two values are compared. If pruning the subtree at node  $N$  would result in a smaller cost complexity, then the subtree is pruned. Otherwise, it is kept.

C4.5 uses a different method called **pessimistic pruning**, which is similar to the cost complexity method in that it also uses error rate estimates to make decisions regarding subtree pruning. Pessimistic pruning, however, does not require the use of a prune set. Instead, it uses the training set to estimate error rates. Recall that an estimate of accuracy or error based on the training set is overly optimistic and, therefore, strongly biased. The pessimistic pruning method therefore adjusts the error rates obtained from the training set by adding a penalty, so as to counter the bias incurred.

## 4 A Modified Decision Tree Algorithm

### 4.1 Results of Classic Algorithms

In this part, we conduct an empirical experiment on the three algorithms proposed. We have selected a dataset on [kaggle.com](https://www.kaggle.com) about car purchases. One key characteristic of the dataset is that the data is arranged *partially in order in terms of certain attributes*. If we don't alter the sequence of the data when splitting the training set and the data set, then there is a *good chance of overfitting* as the training set will not be comprehensive enough. We have selected the dataset

such that all attributes are discrete, and some examples are illustrated in Table 1. The dataset contains 1728 queries, and we have used 70% of the data for training and 30% for testing. During training, we employ 10-fold cross validation to select the parameters for the two algorithms. The implementation is accomplished by the `sklearn` package, and we run the code using the three proposed algorithms respectively with `max_depth` ranging from 3 to 8, `min_sample` ranging from 5 to 50 with interval 5. During the process, we found `min_samples` doesn't significantly influence the result. The test results are shown in Table 2, where `min_sample` are set to 20. Note that here we do not change the sequence of the data to simulate the case of overfitting.

Table 1: Some examples of the selected dataset

Price	Maintenance	Door_number	Capacity	Luggage	Safety	Acceptability
vhigh	vhigh	2	2	small	low	unacc
vhigh	vhigh	2	2	small	med	unacc
vhigh	vhigh	2	2	small	high	unacc
vhigh	vhigh	2	2	med	low	unacc

Table 2: Test results

ALG	Depth	Acc	ALG	Depth	Acc	ALG	Depth	Acc
gini	3	0.902778	entropy	3	0.902778	log_loss	3	0.902778
gini	4	0.819444	entropy	4	0.819444	log_loss	4	0.819444
gini	5	0.861111	entropy	5	0.861111	log_loss	5	0.861111
gini	6	0.861111	entropy	6	0.861111	log_loss	6	0.861111
gini	7	0.861111	entropy	7	0.861111	log_loss	7	0.861111
gini	None	0.851852	entropy	None	0.847222	log_loss	None	0.851852

From the result, we can see that the results for the three respective algorithms are **identical** under three conditions. We deduce that this might be caused by the pruning algorithm, which significantly limits the depths of the tree. This move imposes great constraint and eliminates the potential differences caused by the selection of different metrics. We then select the best result with `max_depth=3` and plot the ROC curve in Figure 2, from which we assert that the binary classification task achieves decent result.




Figure 2: The ROC curve of the best classification result

## 4.2 Modified Pruning based on Dropout

Through the above analysis, it is not difficult to find that if the calculation method of entropy is fixed, the shape of this tree has been roughly determined. However, due to different pruning methods, our final decision tree may still have significant differences. If the quality of the training set is poor, the decision tree we get will have the problem of overfitting. In order to solve the problem of overfitting, we propose the following method.

When considering the solution to this problem, we follow the idea of by *the dropout method* which is commonly used in the construction of neural networks in the field of machine learning. The core idea of the Dropout method can be summarized as follows: different fixed neural networks will have different level of overfitting, and multiple averaging may create some opposite effects. Dropout is a method to force different neuron to become inactive based on a preset probability, which can be regarded as the average of multiple models, similar to the strategy of winning most votes. At the same time, since the classification criterion for each node do not necessarily take effect simultaneously, randomly deleting splitting nodes can reduce the interdependence between features (i.e. classification criteria) and improve the robustness of the decision tree we generate.

Specifically, we introduce a variable  $p$  that represents the possibility of splitting an attribute. It can be regarded as a hyper-parameter, and we hope that the introduction of such a variable can improve the performance during testing. We have conducted an experiment based on idea, and the result is shown in Figure 3. We run the experiment by splitting the dataset into testing set and training set, each accounting for 75% and 25%. We can see that simply introducing  $p$  is not

enough, as the average accuracy drops when  $p$  gradually increases.




Figure 3: Running result of modified decision tree with the introduced parameter  $p$

To deal with the problem, we consider the following. In fact, there is an inheritance relationship in the classification of decision trees. The splitting of the previous node will be directly retained in the classification criteria of the subsequent nodes. In order to eliminate the impact of parameter inheritance, we added a variable  $q$  to adjust the probability of deletion between different levels. As the number of layers increases, the number of preceding nodes of a point also increases, and the probability of using this point as the correct classification criterion also decreases. Therefore, we choose an exponential correction factor to adjust  $p$  dynamically as the level increases. The equation is give by

$$\text{Dropout} = p(1 + q)^l \quad (13)$$

where  $p$  is the initial value of deletion probability, and  $q$  is the correction factor for the number of layers and  $l$  is the level of the decision tree. We then run the test again, this time splitting the original dataset into two parts as before, with 75% for training and 25% for testing. During the testing, we set  $p$  ranging from 0 to 0.25 with interval 0.05, and  $q$  ranging from 0.1 to 1 with interval 0.1. The result can be summarized in Figure 4.

We can see that the classical algorithm achieves an accuracy of 0.86 (also corresponds to the result in Table 2). The best accuracy achieved is 0.96 with  $q > 0.9$ . From Table 2, we can see that this is better than the best result achieved by the three classical algorithms. Multiple hyper-parameter settings also achieved decent performance. But we can also see that there are a noticeable number of parameter combinations that are quite poor. After repetitive testing, we can conclude that the accuracy 0.96 can be achieved with some  $q \in [0.6, 1]$ . This implies that with the

help of adequate pruning, the running result of the algorithm can be improved without changing the attribute selection method.




Figure 4: Running result of modified decision tree with the introduced parameter  $p, q$

### 4.3 Strengths and Limitations

From the previous part, we have proved that the proposed modification to the ID3 algorithm can achieve better result so as to avoid over-fitting. We now try to explain the strength based on information theory.

The ID3 algorithm selects metrics based on maximizing the information gain, which is actually mutual information. In the case where training data is insufficient, the mutual information between certain attributes can be wrongly estimated. The method that we have proposed only keep those with the greatest information gain, and introduce randomness to deal with more trivial nodes. By selecting the right parameters, there stands a better chance of generalizing the model to more unfamiliar test cases. To test the theory, we plot the training accuracy with respect to different  $p$  and  $q$  in Figure 5. We can see that the training accuracy of the classic ID3 algorithm is 1.0, which is a clear sign of overfitting as the test accuracy is only 0.86. On the other hand, the training accuracy of the proposed algorithm mostly floats around 0.85 to 0.95, which implies better flexibility.

However, if the training data is *accurate* enough, then the pruning algorithm can hurt the accuracy, as it potentially eliminates cases where it should be classified with more detail. To test our hypothesis, we run the algorithm with a shuffled dataset, which eliminates overfitting.

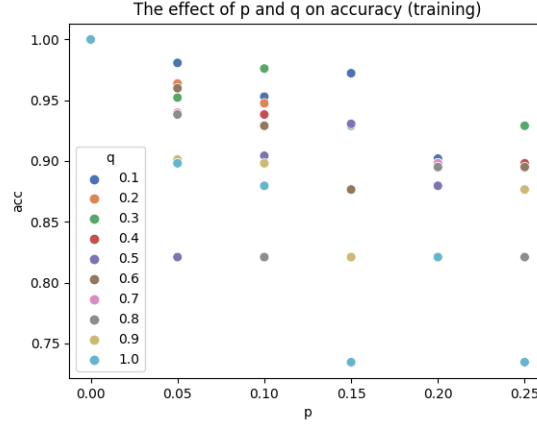


Figure 5: Training accuracy with respect to different  $p$  and  $q$

The results are shown as follows. We can see that as  $p$  increases, the mean of the accuracy decreases, accompanied with increased variance. Nevertheless, the mean value of the accuracy remains acceptable, so setting  $p$  and  $q$  to a relatively low value is robust enough to deal with various test cases.




Figure 6: Experiment Result without overfitting

In conclusion, the introduction of hyper-parameters  $p$  and  $q$  can improve the performance of the decision tree algorithm under certain overfitting circumstances. By adjusting the values of  $p$  and  $q$ , we can find the optimal balance between underfitting and overfitting, and achieve better accuracy on new data. But it lacks stability and can perform worse if the training data is more accurate and well-defined.

## References

- [1] Leo Breiman. *Classification and regression trees*. Routledge, 2017.
- [2] Thomas M Cover. *Elements of information theory*. John Wiley & Sons, 1999.
- [3] What Is Data Mining. Data mining: Concepts and techniques. *Morgan Kaufmann*, 10:559–569, 2006.
- [4] J. Ross Quinlan. Induction of decision trees. *Machine learning*, 1:81–106, 1986.
- [5] J Ross Quinlan. *C4. 5: programs for machine learning*. Elsevier, 2014.
- [6] Claude E Shannon. A mathematical theory of communication. *The Bell system technical journal*, 27(3):379–423, 1948.